

# Computer Science for the Physical Sciences

*Week 4*

---

*Craig Rasmussen (Research Support Services, University of Oregon)*


# Which one is most like your brain?





# Spiking neurons and dendritic trees

---



# Are computers fast enough to compete?


---

- $10^9$  neurons x  $10^3$  Hz x 100 operations
  - $\sim 10^{14}$  op/s
- Roadrunner computer at Los Alamos
  - $16 \times 180 = 2880$  hybrid nodes with 4 IBM cell cores
  - each cell core has 8 SPU cores
    - $2880 \times 4 \times 8 = 92\text{K}$  compute cores
  - each SPU  $\sim 200 \times 10^9$  flops (single precision)
  - $92\text{K} \times 200 \times 10^9 > 10^{15}$  flops



# Are computers big enough to compete?

- Visual cortex
  - $10^9$  neurons
  - $10^{13}$  neurons synaptic connections
- Los Alamos' Roadrunner
  - 92K cores
  - 10K transistors per core
  - $10^9$  transistors
  - where are the equivalent connections (memory pathways)?



# Do computers have enough memory to compete?

---

- Assume memory is in the synaptic connections
  - $10^9$  neurons x  $10^4$  connections x 1 byte
  - $\sim 10^{13}$  bytes
- Roadrunner
  - $> 10^{13}$  bytes

# Do computers have enough bandwidth to compete?

---

- Assume the bandwidth is in the synaptic connections
  - $10^{13}$  connections x 1000 Hz
  - $10^{16}$  bytes/sec
- Roadrunner
  - 11520 Cell processors x 20 Gbytes/sec
  - $10^{14}$  bytes/sec


# So why can't thinking machines think?

---

- Big enough
- Fast enough
- Have enough memory
- But lack bandwidth
  - using memory to simulate synaptic connections
  - must time share access to memory
- So what, we'll just run 100-1000 times slower
- And Voila, we'll have a brain!
- But we don't know the circuit (but can be learned with STDP)




But LANL Roadrunner required 2.35 MW.  
So why doesn't your just brain melt?



# Ivan Sutherland: Should time be quantized?

---

- A chip is big
- A chip is clocked at 2-3 GHz
- Takes several clock cycles to traverse the chip
- Therefore modern chips have many clocks (10K?)
- A transistor does something each clock cycle
  - mostly nothing but waste energy
- Neurons only fire when necessary (mostly true)
- Ivan Sutherland examining computer circuits that are asynchronous
  - only active when necessary



# Computer Science Minor: *Last week*

---

- Required courses (24 credits)
  - Introduction to Computer Science I-II-III
  - Elements of Discrete Mathematics I-II
  - Introduction to Data Structures Lists and Maps
- Upper-division courses (8 credits)
  - Computer Architecture
  - Introduction to Algorithms Complexity
  - C/C++ and Unix Python and Shell
  - Operating Systems
  - Automata Theory
  - Software Methodology I-II  
Revision Control and Make Files
  - Introduction to Compilers
  - Computational Science
  - Bioinformatics
  - Data Mining
  - Introduction to Artificial Intelligence
  - Machine Learning

# Computer Science Minor: *This week*

---

- Required courses (24 credits)
  - Introduction to Computer Science I-II-III
  - Elements of Discrete Mathematics I-II
  - Introduction to Data Structures
- Upper-division courses (8 credits)
  - Computer Architecture
  - Introduction to Algorithms
  - C/C++ and Unix Functions and pipes
  - Operating Systems
  - Automata Theory
  - Software Methodology I-II  
Revision Control and Make Files
  - Introduction to Compilers
  - Computational Science
  - Bioinformatics
  - Data Mining
  - Introduction to Artificial Intelligence
  - Machine Learning

# Computer Science Minor: *Make files*

---

- Required courses (24 credits)
    - Introduction to Computer Science I-II-III
    - Elements of Discrete Mathematics I-II
    - Introduction to Data Structures
  - Upper-division courses (8 credits)
    - Computer Architecture
    - Introduction to Algorithms
    - C/C++ and Unix **Functions and pipes**
    - Operating Systems
    - Automata Theory
    - Software Methodology I-II
    - Introduction to Compilers
    - Computational Science
    - Bioinformatics
    - Data Mining
    - Introduction to Artificial Intelligence
    - Machine Learning
- Revision Control and Make Files**



A *makefile* maintains groups of programs based on dependencies being satisfied

---

```
#
# this is a comment

#
# define environment variables (compilers/linker/libraries...)

CC = gcc

#
# define targets
all: hello
hello.o: hello.c
    $(CC) -c hello.c -o hello.o
hello: hello.o
    $(CC) -o hello hello.o

# run tests
check:

# clean up
clean:
    rm -f hello.o hello
```

target →

dependency →

tab →

# Computer Science Minor: Functions, classes, and Unix pipes

---

- Required courses (24 credits)
    - Introduction to Computer Science I-II-III
    - Elements of Discrete Mathematics I-II
    - Introduction to Data Structures
  - Upper-division courses (8 credits)
    - Computer Architecture
    - Introduction to Algorithms
    - C/C++ and Unix **Functions and pipes**
    - Operating Systems
    - Automata Theory
    - Software Methodology I-II
    - Introduction to Compilers
    - Computational Science
    - Bioinformatics
    - Data Mining
    - Introduction to Artificial Intelligence
    - Machine Learning
- Revision Control and Make Files

# Classes

---

- A class encapsulates functions and state variables
- Class Foo
  - `int x, y, z; // state variables`
  - `void f1(); // function`
- A class is a template (recipe) for creating objects
- A program can have pointers to many live objects at once
- Each object contains state
- In parallel programming **state is evil!**
  - who modified `x` and when?

# Functions

---

- A function takes input and produces output
- Functions are composable
  - $f_3(f_2(f_1(x)))$
- What happens to the state variables?
  - $f_2()$  consumes the output of  $f_1()$
- Going stateless is good!

# Unix pipes

---

- A unix shell program takes input and produces output
  - standard input (file)
  - standard output (file)
- Unix shell programs are composable with pipes
  - `program1 | program2 | program3`
  - the output of program1 is said to be “piped” to the input of program2
- What happens to the state variables (files)?
  - program2 consumes the output of program1



# Computer Science Minor: Revision control

---

- Required courses (24 credits)
    - Introduction to Computer Science I-II-III
    - Elements of Discrete Mathematics I-II
    - Introduction to Data Structures
  - Upper-division courses (8 credits)
    - Computer Architecture
    - Introduction to Algorithms
    - C/C++ and Unix **Functions and pipes**
    - Operating Systems
    - Automata Theory
    - Software Methodology I-II
    - Introduction to Compilers
    - Computational Science
    - Bioinformatics
    - Data Mining
    - Introduction to Artificial Intelligence
    - Machine Learning
- Revision Control and Make Files**

# Git is now the standard version control system


---

- Configuring Git

- `git config --global user.name "Your Name"`
- `git config --global user.email "user@uoregon.edu"`
- `git config --global core.editor emacs`

- Creating a new shared repository


- `mkdir repo`
- `cd repo`
- `git init --bare`



# Creating a local copy of an existing repository and adding files

---


- Clone a repository
  - `git clone /usr/local/repos/repo me`
  - `git clone URL`
- Add a file
  - `cd me`
  - `touch README`
  - `git add README`
  - `git commit -m "Initial version." README`



# Sharing changes to a file

---

- Edit the file then compare changes
  - emacs README
  - git diff README
- Discovery
  - git status
- Commit changes to local repository
  - git commit -m "a message" README
- Push changes to shared repository
  - git push



# Retrieving changes that a teammate has made

---

- Fetch the latest from the shared repository
  - `git fetch`
- Merge with local private repository
  - `git merge origin/master`
- If you have merge conflicts you must fix them
  - `emacs README`
  - `git add README`
  - `git commit -m "Liked my changes better"`

